

LendSmart_Analysis

November 9, 2025

1 LendSmart Credit Risk Analysis

1.0.1 Luis Alan Morales Castillo A01659147

1.0.2 Paulina Díaz Arroyo A010295932

1.0.3 Rodrigo Jiménez Ortiz A01029623

Section 1: Project Setup & Data Loading

- Import all necessary libraries

```
[2]: import pandas as pd
import numpy as np
import sklearn as sk
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import RocCurveDisplay, roc_auc_score
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.preprocessing import StandardScaler
from sklearn.discriminant_analysis import
    LinearDiscriminantAnalysis, QuadraticDiscriminantAnalysis
```

- Load the credit_risk_data.csv file.

```
[3]: df=pd.read_csv("../datos/credit_risk_data-1.csv")
df.head()
```

```
[3]:   application_id application_date  loan_amount  annual_income \
0          APP_2328      2022-01-01     132221.82      60451.82
1          APP_558       2022-01-01     134906.42      114634.08
2          APP_2477      2022-01-01      30285.19      82772.53
3          APP_741       2022-01-01      32516.09      94023.36
4          APP_145       2022-01-02     77900.99      53515.02

  employment_years  job_stability_score  credit_score  credit_utilization \
0                  6.6                 0.898           679            0.106
```

```

1          10.3           0.808        718       0.030
2          12.1           0.964        768       0.174
3          9.1            0.690        670       0.141
4          7.2            0.679        651       0.097

    payment_history_score  open_credit_lines  debt_to_income_ratio \
0                  0.876                 1             0.451
1                  0.719                 4             0.090
2                  0.775                 6             0.201
3                  0.993                 3             0.322
4                  0.946                 2             0.222

    savings_ratio  asset_value  age education_level marital_status \
0          0.500     352569.55   41      High School      Married
1          0.235     224364.21   46          Masters     Divorced
2          0.172     514765.55   44      High School     Widowed
3          0.368     182541.72   26      Bachelors      Single
4          0.324     223691.29   50      Associates      Single

    residential_stability  loan_status
0                  3.5          0
1                  11.4         0
2                  8.6          0
3                  3.9          0
4                  9.6          0

```

- Perform an initial inspection: `.head()`, `.info()`, `.describe()`.

```
[ ]: display(df.head())
display(df.describe())
display(df.info())

    application_id application_date  loan_amount  annual_income \
0      APP_2328      2022-01-01    132221.82    60451.82
1      APP_558       2022-01-01    134906.42    114634.08
2      APP_2477      2022-01-01    30285.19     82772.53
3      APP_741       2022-01-01    32516.09     94023.36
4      APP_145       2022-01-02    77900.99     53515.02

    employment_years  job_stability_score  credit_score  credit_utilization \
0                  6.6                0.898        679          0.106
1                 10.3                0.808        718          0.030
2                 12.1                0.964        768          0.174
3                 9.1                0.690        670          0.141
4                 7.2                0.679        651          0.097

    payment_history_score  open_credit_lines  debt_to_income_ratio \
0                  0.876                 1             0.451

```

1	0.719	4	0.090			
2	0.775	6	0.201			
3	0.993	3	0.322			
4	0.946	2	0.222			
	savings_ratio	asset_value	age	education_level	marital_status	\
0	0.500	352569.55	41	High School	Married	
1	0.235	224364.21	46	Masters	Divorced	
2	0.172	514765.55	44	High School	Widowed	
3	0.368	182541.72	26	Bachelors	Single	
4	0.324	223691.29	50	Associates	Single	
	residential_stability	loan_status				
0	3.5	0				
1	11.4	0				
2	8.6	0				
3	3.9	0				
4	9.6	0				
	loan_amount	annual_income	employment_years	job_stability_score	\	
count	2500.000000	2500.000000	2500.000000	2500.000000		
mean	155716.305344	67707.807596	6.675640	0.634643		
std	149605.357952	27302.931731	3.488021	0.293276		
min	5000.000000	15000.000000	0.000000	0.011000		
25%	42984.517500	47475.317500	4.000000	0.375500		
50%	97054.315000	66963.475000	6.700000	0.752000		
75%	213214.992500	87347.642500	9.300000	0.866000		
max	500000.000000	149929.960000	19.300000	0.999000		
	credit_score	credit_utilization	payment_history_score	\		
count	2500.000000	2500.000000	2500.000000			
mean	681.728400	0.358176	0.740733			
std	88.683309	0.289995	0.285966			
min	334.000000	0.004000	0.029000			
25%	642.750000	0.131000	0.517500			
50%	700.000000	0.246000	0.880500			
75%	743.000000	0.592250	0.956000			
max	850.000000	0.998000	1.000000			
	open_credit_lines	debt_to_income_ratio	savings_ratio	asset_value	\	
count	2500.000000	2500.000000	2500.000000	2500.000000		
mean	3.451600	0.408094	0.320784	175666.741236		
std	2.083793	0.224736	0.192079	182652.568930		
min	0.000000	0.009000	0.000000	550.630000		
25%	2.000000	0.228000	0.161000	49513.082500		
50%	3.000000	0.359000	0.327000	121018.750000		
75%	5.000000	0.565000	0.464000	235513.902500		
max	11.000000	0.979000	0.893000	1000000.000000		

```

      age  residential_stability  loan_status
count  2500.000000             2500.000000  2500.000000
mean   42.045600              6.023200    0.265600
std    12.092395              3.205397    0.441741
min    18.000000              0.000000    0.000000
25%   34.000000              3.600000    0.000000
50%   42.000000              5.900000    0.000000
75%   50.000000              8.400000    1.000000
max   75.000000              16.400000   1.000000

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2500 entries, 0 to 2499
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   application_id  2500 non-null   object  
 1   application_date 2500 non-null   object  
 2   loan_amount      2500 non-null   float64
 3   annual_income    2500 non-null   float64
 4   employment_years 2500 non-null   float64
 5   job_stability_score 2500 non-null   float64
 6   credit_score     2500 non-null   int64  
 7   credit_utilization 2500 non-null   float64
 8   payment_history_score 2500 non-null   float64
 9   open_credit_lines 2500 non-null   int64  
 10  debt_to_income_ratio 2500 non-null   float64
 11  savings_ratio    2500 non-null   float64
 12  asset_value      2500 non-null   float64
 13  age               2500 non-null   int64  
 14  education_level  2500 non-null   object  
 15  marital_status   2500 non-null   object  
 16  residential_stability 2500 non-null   float64
 17  loan_status       2500 non-null   int64  

dtypes: float64(10), int64(4), object(4)
memory usage: 351.7+ KB

```

None

- Write a brief summary of your initial findings

The dataset contains 2,500 loan application records from 2022 to 2024, all without null values. The variables cover financial, demographic, and credit aspects of the applicant. The target variable (loan_status) indicates whether the loan was repaid or defaulted on, with a default rate of 26.56%.


The data shows that loan applicants have highly varied financial profiles. For example, the loan amount requested ranges from \$5,000 to \$500,000, with an average of approximately \$155,000. Annual income also varies considerably, from \$15,000 to nearly \$150,000. Most customers have a credit score of 682, but there are also individuals with very low scores. Indicators such as credit utilization, job stability score, and payment history score reflect that many customers have good

credit history, although there are extreme values that could indicate risk. Furthermore, the debt-to-income ratio shows that some applicants have a high debt burden relative to their income. Variables such as savings ratio, asset value, and age also show significant differences between individuals. All of this indicates that LendSmart needs an analysis model capable of accurately distinguishing between reliable and unreliable customers.

Section 2: Exploratory Data Analysis (EDA)

- **Target Variable:** Plot the distribution of loan_status and calculate the exact default rate.

```
[5]: default_rate=(len(df[df["loan_status"]==1])/len(df))*100
plt.figure()
df["loan_status"].value_counts().plot(kind="bar")
plt.title("Loan Status Distribution")
plt.xticks(rotation=0)
plt.ylabel("Count")
plt.annotate(f'Default Rate: {default_rate:.2f}%', xy=(0.5, 0.9),
             max(df["loan_status"].value_counts())*0.9)),
plt.show()
```




- **Continuous Variables:** For key predictors (credit_score, annual_income, debt_to_income_ratio, etc.), create plots (e.g., box plots or histograms) that com-


pare the distribution for defaulters (1) vs. non-defaulters (0).


```
[84]: numeric_features=["loan_amount", "annual_income",  
    ↴"employment_years", "credit_score", "payment_history_score",  
    ↴"debt_to_income_ratio", "job_stability_score", "credit_utilization",  
    ↴"open_credit_lines", "asset_value", "age", "residential_stability", "savings_ratio"]  
df[numeric_features + ["loan_status"]].corr()["loan_status"].abs().  
    ↴sort_values(ascending=False)
```


```
[84]: loan_status      1.000000  
payment_history_score  0.931093  
job_stability_score   0.912006  
credit_utilization    0.906134  
debt_to_income_ratio  0.785210  
credit_score           0.728160  
savings_ratio          0.697444  
employment_years       0.628269  
residential_stability 0.531898  
annual_income          0.508140  
asset_value            0.430579  
open_credit_lines      0.429446  
age                    0.355555  
loan_amount             0.003192  
Name: loan_status, dtype: float64
```

```
[86]: key_predicors=["payment_history_score", "job_stability_score",  
    ↴"credit_utilization", "debt_to_income_ratio", "credit_score"] # Top 5  
    ↴predictors based on correlation with loan_status  
for predictor in key_predicors:  
    # Boxplots  
    plt.figure()  
    sns.boxplot(x="loan_status", y=predictor, data=df)  
    plt.title(f"{predictor} vs Loan Status")  
    plt.xlabel("0 = Good (paid in full), 1 = Default (90+ days pastdue)")  
    plt.ylabel(predictor)  
    plt.show()  
    # Histograms  
    plt.figure()  
    sns.histplot(data=df, x=predictor, hue="loan_status", kde=True,  
    ↴element="step", stat="count")  
    plt.title(f"{predictor} Distribution by Loan Status")  
    plt.xlabel(predictor)  
    plt.ylabel("Count")  
    plt.show()
```













credit_utilization vs Loan Status









debt_to_income_ratio Distribution by Loan Status







- **Categorical Variables:** Create bar plots showing the mean default rate for each category in education_level and marital_status.

```
[87]: edu_default = df.groupby('education_level')['loan_status'].mean().sort_values()
plt.figure(figsize=(8, 5))
ax = sns.barplot(x=edu_default.index, y=edu_default.values, palette='Blues_d')

plt.title('Default Rate by Education Level', fontsize=14, weight='bold')
plt.xlabel('Education Level', fontsize=12)
plt.ylabel('Default Rate', fontsize=12)
plt.xticks(rotation=45)


for i, rate in enumerate(edu_default.values):
    plt.text(i, rate + 0.01, f'{rate:.2%}', ha='center', va='bottom', fontsize=10)
plt.tight_layout()
plt.show()
```

C:\Users\Asus\AppData\Local\Temp\ipykernel_13240\3893837441.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in

v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
ax = sns.barplot(x=edu_default.index, y=edu_default.values, palette='Blues_d')
```




```
[ ]: marital_default = df.groupby('marital_status')['loan_status'].mean().  
    ↪sort_values()  
  
plt.figure(figsize=(8, 5))  
ax = sns.barplot(x=marital_default.index, y=marital_default.values, ↪  
    ↪palette='Greens_d')  
  
plt.title('Default Rate by Marital Status', fontsize=14, weight='bold')  
plt.xlabel('Marital Status', fontsize=12)  
plt.ylabel('Default Rate', fontsize=12)  
  
for i, rate in enumerate(marital_default.values):  
    plt.text(i, rate + 0.01, f'{rate:.2%}', ha='center', va='bottom', ↪  
    ↪fontsize=10)  
  
plt.tight_layout()  
plt.show()
```

C:\Users\Asus\AppData\Local\Temp\ipykernel_13240\585239893.py:4: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
ax = sns.barplot(x=marital_default.index, y=marital_default.values,  
palette='Greens_d')
```



- **Correlations:** Generate a correlation matrix heatmap for all numeric predictors.

```
[ ]: plt.figure(figsize=(12, 8))  
sns.heatmap(df.corr(numeric_only=True), annot=True, fmt=".2f", cmap='inferno')  
plt.title('Correlation Matrix')  
plt.tight_layout()  
plt.show()
```



Section 3: Data Preprocessing

- **Handle Categorical Data:** Convert education_level and marital_status into numerical dummy variables (pd.get_dummies()).

```
[90]: dummies_edu = pd.get_dummies(df["education_level"], prefix="edu", ↴
    drop_first=True)
dummies_mar = pd.get_dummies(df["marital_status"], prefix="marital", ↴
    drop_first=True)

dummies_edu
```

```
[90]:      edu_Bachelors  edu_Doctorate  edu_High School  edu_Masters
0           False        False          True        False
1           False        False         False        True
2           False        False          True        False
3            True        False         False        False
4           False        False          False        False
...
2495          True        False          False        False
2496          False        False         False        False
2497          False        False         False        True
2498          False        False          True        False
```

```
2499      True     False     False     False
```

```
[2500 rows x 4 columns]
```

- **Define Predictors (X) and Target (y):** Create X (features) and y (target) variables. Be sure to drop non-predictive columns like application_id.

```
[91]: x_features=["credit_score", "payment_history_score", "debt_to_income_ratio",  
                 "job_stability_score", "credit_utilization"]  
X = pd.concat([df[x_features], dummies_edu, dummies_mar], axis=1)  
  
y_target=df["loan_status"]
```

- **Train-Test Split:** Split the data into X_train, X_test, y_train, y_test. Use test_size=0.2 and a random_state=42 for reproducibility.

```
[92]: x_train, x_test, y_train, y_test=train_test_split(X, y_target, test_size=0.2,  
                                                 random_state=42, stratify=y_target)
```

- **Standardization:** Initialize a StandardScaler. **Fit** it only on X_train, then **transform** both X_train and X_test. This is critical for interpreting LDA coefficients.

```
[ ]: scaler = StandardScaler()  
x_train_scaled = scaler.fit_transform(x_train)  
x_test_scaled = scaler.transform(x_test)
```

Section 4: Statistical Assumption Testing

- **Key Statistical Assumptions Differentiating LDA and QDA**

Both Linear Discriminant Analysis (LDA) and Quadratic Discriminant Analysis (QDA) assume that the predictors within each class follow a multivariate normal distribution. The main difference lies in how they treat the covariance matrices of these distributions.

- Multivariate Normality:

This assumption means that the predictors for each class should be approximately normally distributed. Based on the EDA plots, most variables (credit_score, credit_utilization, debt_to_income_ratio, job_stability_score) show roughly bell-shaped distributions, though some exhibit moderate skewness (payment_history_score being the most skewed). While not perfectly normal, the data are reasonably smooth and unimodal within each loan status group, suggesting that the multivariate normality assumption is reasonably met for applying discriminant analysis.

- Homogeneity of Covariance Matrices:

LDA assumes that all classes share a single covariance matrix, producing linear decision boundaries, while QDA allows each class to have its own covariance matrix, resulting in quadratic boundaries. If the class scatterplots or spreads differ noticeably—as suggested by some variation in variable ranges between “Default” and “Paid in Full” loans—this indicates unequal covariance structures.

- **Hypothesis:** If the covariance matrices are unequal, we expect QDA to outperform LDA, as QDA can better model differing class variances and correlations, though at the cost of greater model complexity.

Section 5: Model 1 - Linear Discriminant Analysis (LDA)

- Initialize LinearDiscriminantAnalysis and fit it on your standardized X_train and y_train.

```
[95]: LDA=LinearDiscriminantAnalysis()
LDA.fit(x_train_scaled, y_train.values.ravel())
```

```
[95]: LinearDiscriminantAnalysis()
```

- Interpret Coefficients: Extract the lda.coef_. Place them in a DataFrame with their feature names. Sort by absolute value.

```
[96]: coef=LDA.coef_[0]
feature_names=x_train.columns

df_coef=pd.DataFrame({"Feature":feature_names, "Coefficient":coef})
df_coef[["Abs_Coefficient"]]=df_coef[["Coefficient"]].abs()
df_coef.sort_values(by="Abs_Coefficient", ascending=False, inplace=True)
df_coef
```

	Feature	Coefficient	Abs_Coefficient
1	payment_history_score	-15.543899	15.543899
3	job_stability_score	-13.070033	13.070033
4	credit_utilization	11.643710	11.643710
2	debt_to_income_ratio	4.506771	4.506771
0	credit_score	-3.937238	3.937238
7	edu_High_School	1.133565	1.133565
10	marital_Single	0.515575	0.515575
9	marital_Married	-0.492058	0.492058
6	edu_Documentary	-0.417348	0.417348
5	edu_Bachelors	-0.181114	0.181114
8	edu_Masters	-0.174104	0.174104
11	marital_Widowed	0.051486	0.051486

- Write a clear interpretation: Which 3-5 variables are the most important drivers of default risk? What does the sign (+/-) tell you?

The top predictors of loan default risk, based on the absolute value of LDA coefficients, are:

- Payment History Score (-15.54): A poor payment history greatly increases default risk. Customers who have missed prior payments or shown inconsistent repayment patterns are significantly more likely to default again.
- Job Stability Score (-13.07): Applicants with unstable employment—frequent job changes or short tenure—tend to face higher financial uncertainty, which increases the likelihood of falling behind on loan payments.

- Credit Utilization (+11.64): High credit utilization indicates that the borrower is already heavily relying on credit. This pressure raises the probability of default, as the borrower has limited capacity to absorb additional debt.
- Debt-to-Income Ratio (+4.51): This is one of the strongest predictors. Applicants whose debt obligations take up a large share of their income are much more likely to struggle with repayment and default.
- Credit Score (-3.94): Lower credit scores strongly correlate with higher default risk. Credit Score remains a reliable indicator of financial discipline and borrower reliability.

Section 6: Model 2 - Quadratic Discriminant Analysis (QDA)

- Initialize QuadraticDiscriminantAnalysis and fit it on your standardized X_train and y_train

```
[97]: QDA = QuadraticDiscriminantAnalysis()
QDA.fit(x_train_scaled, y_train)
```

```
[97]: QuadraticDiscriminantAnalysis()
```

Section 7: Model Evaluation & Comparison

- Generate predictions (.predict()) for both models on the X_test data.

```
[98]: y_pred_lda = LDA.predict(x_test_scaled)
y_pred_qda = QDA.predict(x_test_scaled)
```

- Confusion Matrices: For both models, generate and plot a ConfusionMatrixDisplay.


```
[99]: def plot_confusion_matrix(y_true, y_pred, model_name):
    cm = confusion_matrix(y_true, y_pred)
    labels = ["No Default (0)", "Default (1)"]


    plt.figure(figsize=(6, 5))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,
                xticklabels=labels, yticklabels=labels, linewidths=0.5, linecolor='gray')

    plt.title(f"Confusion Matrix - {model_name}", fontsize=14, weight='bold')
    plt.xlabel("Predicted Label", fontsize=12)
    plt.ylabel("True Label", fontsize=12)
    plt.xticks(rotation=0)
    plt.yticks(rotation=0)
    plt.tight_layout()
    plt.show()

plot_confusion_matrix(y_test, y_pred_lda, "LDA")
plot_confusion_matrix(y_test, y_pred_qda, "QDA")
```

Confusion Matrix - LDA





- Classification Reports: For both models, print the classification_report.

```
[100]: print("Classification Report - LDA")
print(classification_report(y_test, y_pred_lda))
```

```
print("Classification Report - QDA")
print(classification_report(y_test, y_pred_qda))
```

Classification Report - LDA				
	precision	recall	f1-score	support
0	1.00	1.00	1.00	367
1	1.00	1.00	1.00	133
accuracy			1.00	500
macro avg	1.00	1.00	1.00	500
weighted avg	1.00	1.00	1.00	500

Classification Report - QDA

	precision	recall	f1-score	support
0	1.00	1.00	1.00	367
1	1.00	1.00	1.00	133
accuracy			1.00	500
macro avg	1.00	1.00	1.00	500
weighted avg	1.00	1.00	1.00	500

- ROC Curves: Generate the RocCurveDisplay for both models and plot them on the same axis for a direct visual comparison. Report the AUC (Area Under the Curve) score for both.

```
[80]: y_proba_lda = LDA.predict_proba(x_test_scaled)[:, 1]
y_proba_qda = QDA.predict_proba(x_test_scaled)[:, 1]

plt.figure(figsize=(8, 6))
RocCurveDisplay.from_predictions(y_test, y_proba_lda, name="LDA", color="blue")
plt.plot([0, 1], [0, 1], '--', color='gray')
plt.title("ROC Curve - LDA")
plt.grid(True)
plt.show()

plt.figure(figsize=(8, 6))
RocCurveDisplay.from_predictions(y_test, y_proba_qda, name="QDA", color="orange")
plt.plot([0, 1], [0, 1], '--', color='gray')
plt.title("ROC Curve - QDA")
plt.grid(True)
plt.show()

# Combined plot
fig, ax = plt.subplots(figsize=(8, 6))


RocCurveDisplay.from_predictions(y_test, y_proba_lda, name="LDA", ax=ax)
RocCurveDisplay.from_predictions(y_test, y_proba_qda, name="QDA", ax=ax)

ax.plot([0, 1], [0, 1], linestyle='--', color='gray')
ax.set_title("ROC Curve - LDA vs QDA")
ax.grid(True)


plt.tight_layout()
plt.show()
auc_lda = roc_auc_score(y_test, y_proba_lda)
auc_qda = roc_auc_score(y_test, y_proba_qda)


print(f"AUC Score - LDA: {auc_lda:.4f}")
print(f"AUC Score - QDA: {auc_qda:.4f}")
```

<Figure size 800x600 with 0 Axes>



<Figure size 800x600 with 0 Axes>





AUC Score - LDA: 1.0000

AUC Score - QDA: 1.0000

Section 8: Technical Conclusion & Model Selection

Based on the evaluation results in Section 7, both the LDA and QDA models achieve perfect classification performance on this test set, as shown by the ROC curves for both models reaching the top-left corner of the chart and yielding an AUC of 1.0. This means that each model correctly identified every single default and non-default case, with no false positives and no false negatives. In other words, the predicted probabilities for defaulters were always higher than those for non-defaulters, showing complete separability between the two groups. Although this kind of perfect performance is extremely rare in real-world lending environments, it is statistically possible when the underlying predictors are highly informative and the dataset exhibits a strong, clean distinction between risk profiles. The overlap of the two curves indicates that LDA and QDA produced the same

correct classifications despite using different types of decision boundaries, reinforcing the strength of the financial indicators used. For the business, this result implies that the model can accurately distinguish between high-risk and low-risk borrowers in this sample, providing a powerful tool for reducing default losses with minimal risk of unnecessarily rejecting good customers, while noting that further validation on real-world data is recommended to confirm this level of performance.

LDA works because it assumes that both classes share the same covariance structure, which creates a clean linear boundary between defaulters and non-defaulters. This makes it stable, easy to interpret, and highly reliable when the predictors are well-separated. QDA works because it allows each class to have its own covariance matrix, giving it the flexibility to model more complex, curved decision boundaries. When the data is highly structured and the class distributions differ significantly, QDA can capture these patterns and achieve perfect separation.

In this case, either model works equally well, but LDA may be preferred for its simplicity and interpretability, especially since the results were the same. This means that LDA got the same results as QDA, but with simpler assumptions.